

KESTREL HOME APPLIANCES — MEMORANDUM

TO: Ritu Deshpande, Head of D2C Operations
CC: Farhan Sheikh, Meenal Joshi, Tanmay Kulkarni
RE: **Warranty fraud flag — what to do and what to expect**

DATE: 6 October 2026
FROM: Kabir Nanda, Account Lead

The decision. Use the tool to pick which claims the investigation desk looks at each month, and stop judging it on accuracy. Accuracy cannot tell a useful tool from a useless one: in June, paying every claim without checking would have been 96.9% “accurate” — almost the board’s 97% — while stopping no fraud. A tool that actually holds suspicious claims scores about 95%, because some genuine claims get held too.

The number that matters: rupees per claim checked. Replaying the past (training only on older claims, then scoring June), the desk’s 40 reviews would have found about 12 of the 22 fraudulent claims, stopping about Rs 18,000. It would also have held 26 genuine customers, costing about Rs 9,900 in goodwill (Rs 380 each). Net: about **Rs 220 for every claim checked**. A second replay gave about Rs 130. Plan on **Rs 130–220 per check, roughly Rs 5,000–9,000 a month**, catching roughly half of the fraud. These are small samples (22 fraud cases in June), so treat them as a range. This is a modest, real gain, not a silver bullet.

What the data says about the new partners. You were half right. Fraud has risen sharply since the 1 May change (claims under Rs 2,000 are paid without inspection): new partners went from no fraud before May to 16% of claims since. But it is concentrated in **seven outlets**: SP3160, SP3232, SP3318, SP3129, SP3319, SP3118 and SP3286. The other roughly 38 new partners with claims since May had 1 fraud case in 153 claims — Meenal is right that they are fine. Two of the seven (SP3118, SP3286) rest on only two claims each, so treat them as strong suspicion rather than proof.

What to do next week

- 1. Audit the seven outlets** and take back their auto-approve right (inspect their small claims) until cleared. About 80% of what the tool flags for July–September sits on these outlets, so this does more than 120 separate checks.
- 2. Run the tool on all other claims** and give the desk the top 40 each month. The screen shows the reasons in plain English for each.
- 3. Change the board KPI** to “rupees of fraud stopped per claim checked” with “share of fraud caught” beside it. Accuracy can stay as a reference line.
- 4. Close the open cases.** About 200 claims have no investigation outcome, and old records may hide more. Without outcomes the tool cannot learn.
- 5. Retrain every month.** The fraud pattern changed once already, and a tool that had not seen the new pattern did worse than guessing. It also knows nothing about outlets onboarded in July or later.

What I am unsure about. Outcomes only run to 30 June, and the tool is being used on July–September. If new outlets start fraud at new places, it will miss them at first. Fraud on larger, inspected claims is rarer now, but it was the bigger-rupee fraud before May; it is still scored, just not the focus.